

Analyze_ab_test_results_notebook

June 7, 2021

0.1 Analyze A/B Test Results Project

0.2 Table of Contents

- Section ??
- Section ??
- Section ??
- Section ??

Introduction

A/B tests are very commonly performed by data analysts and data scientists. It is important that you get some practice working with the difficulties of these. For this project, you will be working to understand the results of an A/B test run by an e-commerce website. Your goal is to work through this notebook to help the company understand if they should implement the new page, keep the old page, or perhaps run the experiment longer to make their decision.

Part I - Probability

```
In [1]: import pandas as pd
import numpy as np
import random
import matplotlib.pyplot as plt
%matplotlib inline
#We are setting the seed to assure you get the same answers on quizzes as we set up
random.seed(42)
```

1. Now, read in the `ab_data.csv` data. Store it in `df`.

a. Read in the dataset and take a look at the top few rows here:

```
In [2]: df = pd.read_csv('ab_data.csv')
df.head(10)
```

```
Out[2]:
```

	user_id	timestamp	group	landing_page	converted
0	851104	2017-01-21 22:11:48.556739	control	old_page	0
1	804228	2017-01-12 08:01:45.159739	control	old_page	0
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0
4	864975	2017-01-21 01:52:26.210827	control	old_page	1
5	936923	2017-01-10 15:20:49.083499	control	old_page	0

6	679687	2017-01-19 03:26:46.940749	treatment	new_page	1
7	719014	2017-01-17 01:48:29.539573	control	old_page	0
8	817355	2017-01-04 17:58:08.979471	treatment	new_page	1
9	839785	2017-01-15 18:11:06.610965	treatment	new_page	1

b. Use the cell below to find the number of rows in the dataset.

```
In [3]: df.shape[0]
```

```
Out[3]: 294478
```

c. The number of unique users in the dataset.

```
In [4]: df.nunique()
```

```
Out[4]: user_id      290584
        timestamp    294478
        group         2
        landing_page  2
        converted     2
        dtype: int64
```

d. The proportion of users converted.

```
In [5]: df_converted= len(df.query('converted ==1 '))/len(df['converted'])
        print(df_converted)
```

```
0.11965919355605512
```

```
In [6]: #df.query('converted ==1 & landing_page == "new_page").count()
```

e. The number of times the new_page and treatment don't match.

```
In [7]: new_page_n = df.query('group != "treatment" & landing_page == "new_page"')
        treatment_n = df.query('group == "treatment" & landing_page != "new_page"')
        no_times= len(new_page_n)+ len(treatment_n)
        print(no_times)
```

```
3893
```

f. Do any of the rows have missing values?

```
In [8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 294478 entries, 0 to 294477
Data columns (total 5 columns):
user_id      294478 non-null int64
timestamp    294478 non-null object
```

```

group          294478 non-null object
landing_page   294478 non-null object
converted      294478 non-null int64
dtypes: int64(2), object(3)
memory usage: 11.2+ MB

```

2. For the rows where **treatment** does not match with **new_page** or **control** does not match with **old_page**, we cannot be sure if this row truly received the new or old page. Store your new dataframe in **df2**.

```

In [9]: clean_df=df
        clean_df = df[((df.group=='treatment') & (df.landing_page=='new_page')) | ((df.group=='c
        df2=clean_df

```

```

In [10]: # Double Check all of the correct rows were removed - this should be 0
         df2[((df2['group'] == 'treatment') == (df2['landing_page'] == 'new_page')) == False].sh

```

```

Out[10]: 0

```

a. How many unique **user_ids** are in **df2**?

```

In [11]: df2['user_id'].nunique()

```

```

Out[11]: 290584

```

b. There is one **user_id** repeated in **df2**. What is it?

```

In [12]: len(df2['user_id'].duplicated())

```

```

Out[12]: 290585

```

c. What is the row information for the repeat **user_id**?

```

In [13]: df2[df2.duplicated(['user_id'],keep=False)]

```

```

Out[13]:
   user_id  timestamp  group landing_page  converted
1899  773192  2017-01-09 05:37:58.781806  treatment    new_page         0
2893  773192  2017-01-14 02:55:59.590927  treatment    new_page         0

```

d. Remove **one** of the rows with a duplicate **user_id**, but keep your dataframe as **df2**.

```

In [14]: df2=df2.drop_duplicates(subset=['user_id'], keep='first')

```

```

In [15]: len(df2.duplicated())

```

```

Out[15]: 290584

```

a. What is the probability of an individual converting regardless of the page they receive?

```

In [16]: df_convert= len(df2.query('converted ==1 '))/len(df2['converted'])
         print(df_convert)

```

0.11959708724499628

b. Given that an individual was in the control group, what is the probability they converted?

```
In [17]: control_group = len(df2.query('group=="control" and converted==1'))/len(df2.query('group=="control"'))
        print(control_group)
```

0.1203863045004612

c. Given that an individual was in the treatment group, what is the probability they converted?

```
In [18]: treatment_group = len(df2.query('group=="treatment" and converted==1'))/len(df2.query('group=="treatment"'))
        print(treatment_group)
```

0.11880806551510564

d. What is the probability that an individual received the new page?

```
In [19]: received_new = len(df2.query('landing_page=="new_page"))/len(df2['landing_page'])
        print(received_new)
```

0.5000619442226688

e. Is there sufficient evidence to conclude that the new treatment page leads to more conversions?

I can't say that there is enough evidence to conclude that there is more conversion in new treatment page. As we see in the previous Statistics not all of users has received the new page only half of them and the percent of converted users that received is very close about 12% control group. While about 11% treatment. So overall the conversion rate did not increase but it's a very small difference.

Part II - A/B Test

1. For now, assume that the old page is better unless the new page proves to be definitely better at a Type I error rate of 5%, what should your null and alternative hypotheses be?

p_{old} and p_{new} , which are the converted rates for the old and new pages.

H0: $p_{new} - p_{old} \leq 0$

H1: $p_{new} - p_{old} > 0$

2. Assume under the null hypothesis, p_{new} and p_{old} both have "true" success rates equal to the **converted** success rate regardless of page - that is p_{new} and p_{old} are equal. Furthermore, assume they are equal to the **converted** rate in **ab_data.csv** regardless of the page.

-Perform the sampling distribution for the difference in **converted** between the two pages over 10,000 iterations of calculating an estimate from the null.

a. What is the **conversion rate** for p_{new} under the null?

```
In [20]: df2.head()
```

```
Out[20]:
```

	user_id	timestamp	group	landing_page	converted
0	851104	2017-01-21 22:11:48.556739	control	old_page	0
1	804228	2017-01-12 08:01:45.159739	control	old_page	0
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0
4	864975	2017-01-21 01:52:26.210827	control	old_page	1

```
In [21]: pnew = len(df.query('converted ==1 '))/len(df['converted'])
          print(pnew)
```

```
0.11965919355605512
```

b. What is the **conversion rate** for p_{old} under the null?

```
In [22]: pold = len(df.query('converted ==1 '))/len(df['converted'])
          print(pold)
```

```
0.11965919355605512
```

c. What is n_{new} , the number of individuals in the treatment group?

```
In [23]: no_treatment = len(df2.query('group=="treatment"'))
          print(no_treatment)
```

```
145310
```

d. What is n_{old} , the number of individuals in the control group?

```
In [24]: no_control = len(df2.query('group=="control"'))
          print(no_control)
```

```
145274
```

e. Simulate n_{new} transactions with a conversion rate of p_{new} under the null. Store these n_{new} 1's and 0's in **new_page_converted**.

```
In [25]: new_page_converted = np.random.choice([0, 1], no_treatment, p = [pnew, 1-pnew])
          print(new_page_converted)
```

```
[1 0 1 ..., 1 1 1]
```

f. Simulate n_{old} transactions with a conversion rate of p_{old} under the null. Store these n_{old} 1's and 0's in **old_page_converted**.

```
In [26]: old_page_converted = np.random.choice([0, 1], no_control, p = [pold, 1-pold])
          print(old_page_converted)
```

```
[1 1 1 ..., 1 1 1]
```

g. Find $p_{new} - p_{old}$ for your simulated values from part (e) and (f).

```
In [27]: p_diff = new_page_converted.mean()-old_page_converted.mean()
         print(p_diff)
```

```
0.000215467771097
```

```
In [28]: pvalue_diff = pnew-pold
         print(pvalue_diff)
```

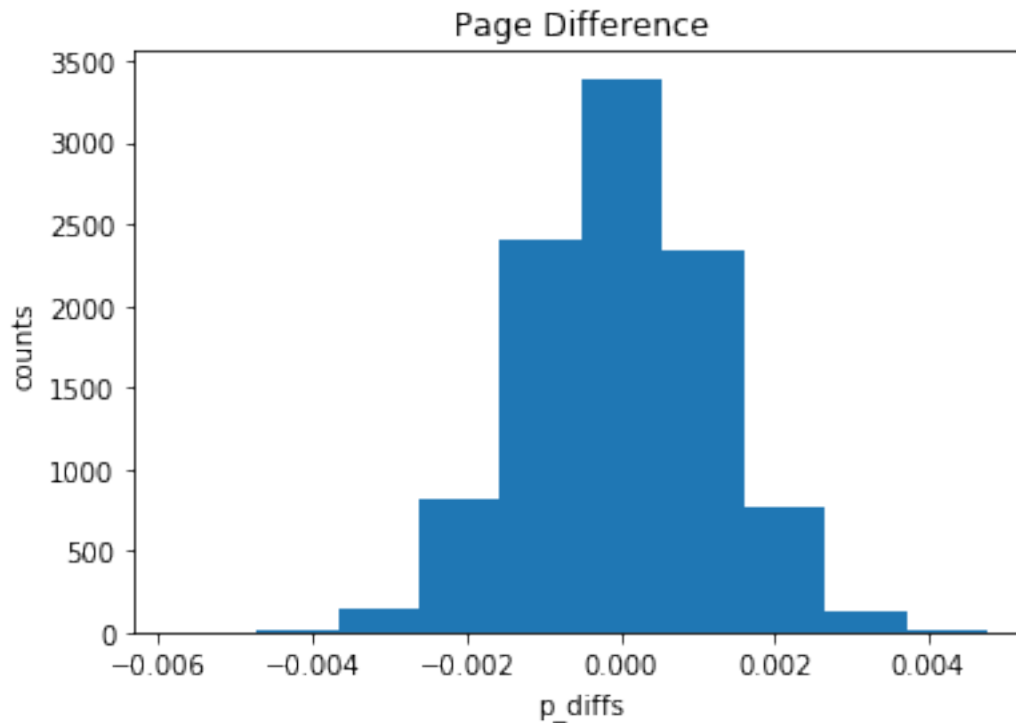
```
0.0
```

h. Create 10,000 $p_{new} - p_{old}$ values using the same simulation process you used in parts (a) through (g) above. Store all 10,000 values in a NumPy array called **p_diffs**.

```
In [29]: p_diffs=[]
         for _ in range(10000):
             new= np.random.choice([0,1],no_treatment, replace=True, p= [pnew, 1-pnew])
             old= np.random.choice([0,1],no_control, replace=True, p= [pold, 1-pold])
             new_mean= new.mean()
             old_mean= old.mean()
             p_diffs.append(new_mean - old_mean)
```

i. Plot a histogram of the **p_diffs**. Does this plot look like what you expected? Use the matching problem in the classroom to assure you fully understand what was computed here.

```
In [30]: p_diffs = np.array(p_diffs)
         plt.hist(p_diffs);
         plt.title ('Page Difference');
         plt.xlabel('p_diffs');
         plt.ylabel('counts');
```

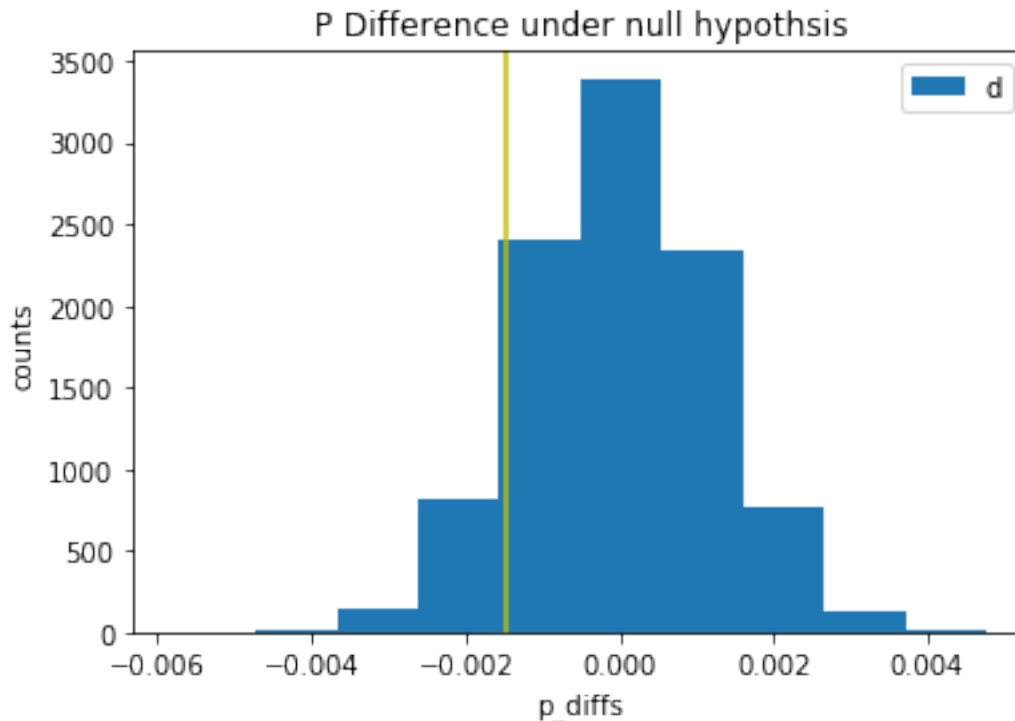


j. What proportion of the **p_diffs** are greater than the actual difference observed in **ab_data.csv**?

```
In [31]: ab_new= df.query('group == "treatment"').converted.mean()
         ab_old= df.query('group == "control"').converted.mean()
         ab_diff= ab_new - ab_old
         (p_diffs > ab_diff).mean()
```

```
Out[31]: 0.88929999999999998
```

```
In [32]: plt.hist(p_diffs);
         plt.title ('P Difference under null hypothesis');
         plt.xlabel('p_diffs');
         plt.ylabel('counts');
         plt.legend("difference between old and new");
         plt.axvline(x=ab_diff, color='y');
```



What is this value called in scientific studies? What does this value mean in terms of whether or not there is a difference between the new and old pages?

The proportion of p-difference is about 89% .p-value is large, so the population is above difference which mean that the new-page is not doing better with comparison to the old one, the result is it would be better if we stick to the null hypothesis.

1. below calculation the number of conversions for each page, as well as the number of individuals who received each page. Let `n_old` and `n_new` refer that the number of rows associated with the old page and new pages, respectively.

```
In [33]: import statsmodels.api as sm
         from scipy.stats import norm
         convert_old = len(df2.query('landing_page == "old_page" & converted == 1'))
         convert_new = len(df2.query('landing_page == "new_page" & converted == 1'))
         n_old = len(df2.query('landing_page == "old_page"'))
         n_new = len(df2.query('landing_page == "new_page"'))
```

```
/opt/conda/lib/python3.6/site-packages/statsmodels/compat/pandas.py:56: FutureWarning: The pandas
from pandas.core import datetools
```

```
In [34]: print('n_old',n_old)
         print('n_new',n_new)
```

```
n_old 145274
n_new 145310
```



```
In [35]: print('convert_old',convert_old)
         print('convert_new',convert_new)
```

```
convert_old 17489
convert_new 17264
```

```
In [36]: z_score, p_value = sm.stats.proportions_ztest([convert_old,convert_new], [n_old, n_new])
         print(z_score,p_value)
```

```
1.31092419842 0.905058312759
```

- n. What do the z-score and p-value you computed in the previous question mean for the conversion rates of the old and new pages? Do they agree with the findings in parts j. and k.?

```
In [37]: print(norm.cdf(z_score))
```

```
0.905058312759
```

```
In [38]: print(norm.ppf(1-(0.05)))
```

```
1.64485362695
```

It's mean null hypothesis of old page conversion rate is greater than the new page conversion. As the p-value is 0.9 and z-score equal to 1.3 (should be equal or more 1.644) less than critical value of 95% confidence, so as result we can't reject null hypothesis.

Part III - A regression approach

1. In this final part, you will see that the result you achieved in the A/B test in Part II above can also be achieved by performing regression.

- a. Since each row is either a conversion or no conversion, what type of regression should you be performing in this case?

Since we have a 2 variable whether converted or not so Logistic Regression is the suitable regression type.

- b. The goal is to use **statsmodels** to fit the regression model you specified in part a. to see if there is a significant difference in conversion based on which page a customer receives.

```
In [39]: df2['intercept']=1
         df2[['control', 'treatment']] = pd.get_dummies(df2['group'])
         df2.head()
```

```
/opt/conda/lib/python3.6/site-packages/ipykernel_launcher.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>

```
"""Entry point for launching an IPython kernel.
/opt/conda/lib/python3.6/site-packages/pandas/core/frame.py:3140: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#>

```
self[k1] = value[k2]
```

```
Out[39]:
```

	user_id	timestamp	group	landing_page	converted	\
0	851104	2017-01-21 22:11:48.556739	control	old_page	0	
1	804228	2017-01-12 08:01:45.159739	control	old_page	0	
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0	
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0	
4	864975	2017-01-21 01:52:26.210827	control	old_page	1	

	intercept	control	treatment
0	1	1	0
1	1	1	0
2	1	0	1
3	1	0	1
4	1	1	0

- c. Use **statsmodels** to instantiate your regression model on the two columns you created in part b., then fit the model using the two columns you created in part b. to predict whether or not an individual converts.

```
In [40]: model = sm.Logit(df2['converted'],df2[['intercept','treatment']])
```

- d. Provide the summary of your model below, and use it as necessary to answer the following questions.

```
In [41]: results = model.fit()
```

```
Optimization terminated successfully.
Current function value: 0.366118
Iterations 6
```

```
In [42]: results.summary2()
```

```
Out[42]: <class 'statsmodels.iolib.summary2.Summary'>
"""
Results: Logit
=====
Model:                Logit                No. Iterations:    6.0000
Dependent Variable:    converted              Pseudo R-squared:  0.000
```

```

Date:                2021-06-07 00:44 AIC:                212780.3502
No. Observations:    290584          BIC:                212801.5095
Df Model:            1              Log-Likelihood:      -1.0639e+05
Df Residuals:        290582          LL-Null:           -1.0639e+05
Converged:           1.0000          Scale:             1.0000

```

```

-----
              Coef.   Std.Err.      z      P>|z|      [0.025   0.975]
-----
intercept    -1.9888    0.0081   -246.6690  0.0000   -2.0046   -1.9730
treatment    -0.0150    0.0114   -1.3109   0.1899   -0.0374    0.0074
=====

```

"""

- e. What is the p-value associated with **ab_page**? Why does it differ from the value you found in **Part II**? **Hint**: What are the null and alternative hypotheses associated with your regression model, and how do they compare to the null and alternative hypotheses in **Part II**?

Z,P-value = -1.3109 , 0.1899

Regression hypotheses is :

H0: $p_{new} = p_{old}$

H1: $p_{new} \neq p_{old}$

while the hypotheses in part 2 is:

H0: $- \leq 0$

H1: $- > 0$

The p-value is different from part 2. In part II the p-value is 0.9. This may be due to the regression model test that put an intercept.

- f. Now, you are considering other things that might influence whether or not an individual converts. Discuss why it is a good idea to consider other factors to add into your regression model. Are there any disadvantages to adding additional terms into your regression model?

I think it's good to consider other factors too because this may help to know other things could affect conversion rate. On the other side I think will be a little bit complex and that could be the disadvantages.

- g. Now along with testing if the conversion rate changes for different pages, also add an effect based on which country a user lives in. I will need to read in the **countries.csv** dataset and merge together your datasets on the appropriate rows.

```

In [43]: countries_df = pd.read_csv('countries.csv')
         df_cont = countries_df.set_index('user_id').join(df2.set_index('user_id'), how='inner')
         df_cont.head()

```

```

Out[43]:
   country  timestamp  group landing_page \
user_id
834778    UK  2017-01-14 23:08:43.304998  control    old_page
928468    US  2017-01-23 14:44:16.387854  treatment    new_page
822059    UK  2017-01-16 14:04:14.719771  treatment    new_page

```

711597	UK	2017-01-22 03:14:24.763511	control	old_page
710616	UK	2017-01-16 13:14:44.000513	treatment	new_page

	converted	intercept	control	treatment
user_id				
834778	0	1	1	0
928468	0	1	0	1
822059	1	1	0	1
711597	0	1	1	0
710616	0	1	0	1

In [44]: df_cont.country.unique()

Out[44]: array(['UK', 'US', 'CA'], dtype=object)

In [45]: df_cont.describe()

	converted	intercept	control	treatment
count	290584.000000	290584.0	290584.000000	290584.000000
mean	0.119597	1.0	0.499938	0.500062
std	0.324490	0.0	0.500001	0.500001
min	0.000000	1.0	0.000000	0.000000
25%	0.000000	1.0	0.000000	0.000000
50%	0.000000	1.0	0.000000	1.000000
75%	0.000000	1.0	1.000000	1.000000
max	1.000000	1.0	1.000000	1.000000

In [46]: df_cont.groupby(['country']).mean()

	converted	intercept	control	treatment
country				
CA	0.115318	1.0	0.496448	0.503552
UK	0.120594	1.0	0.501753	0.498247
US	0.119547	1.0	0.499541	0.500459

In [47]: df_cont[['CA', 'UK', 'US']] = pd.get_dummies(df_cont['country'])

In [48]: df_cont.head()

	country	timestamp	group	landing_page	\
user_id					
834778	UK	2017-01-14 23:08:43.304998	control	old_page	
928468	US	2017-01-23 14:44:16.387854	treatment	new_page	
822059	UK	2017-01-16 14:04:14.719771	treatment	new_page	
711597	UK	2017-01-22 03:14:24.763511	control	old_page	
710616	UK	2017-01-16 13:14:44.000513	treatment	new_page	

	converted	intercept	control	treatment	CA	UK	US
user_id							

834778	0	1	1	0	0	1	0
928468	0	1	0	1	0	0	1
822059	1	1	0	1	0	1	0
711597	0	1	1	0	0	1	0
710616	0	1	0	1	0	1	0

```
In [53]: #Create intercept column for the data set df_new
df_cont['intercept'] = 1

#Fit the model, CA as baseline so not included in lm
lm = sm.Logit(df_cont['converted'], df_cont[['intercept', 'US', 'UK']])
results = lm.fit()
results.summary2()
```

Optimization terminated successfully.
 Current function value: 0.366116
 Iterations 6

```
Out[53]: <class 'statsmodels.iolib.summary2.Summary'>
"""
                                Results: Logit
=====
Model:                        Logit                No. Iterations:    6.0000
Dependent Variable: converted      Pseudo R-squared: 0.000
Date:      2021-06-07 00:45  AIC:                        212780.8333
No. Observations:  290584      BIC:                        212812.5723
Df Model:          2           Log-Likelihood:    -1.0639e+05
Df Residuals:      290581      LL-Null:        -1.0639e+05
Converged:         1.0000      Scale:          1.0000
-----
                                Coef.   Std.Err.    z      P>|z|    [0.025   0.975]
-----
intercept    -2.0375    0.0260   -78.3639  0.0000   -2.0885   -1.9866
US            0.0408    0.0269    1.5178  0.1291   -0.0119    0.0935
UK            0.0507    0.0284    1.7863  0.0740   -0.0049    0.1064
=====
"""
```

- h. Though you have now looked at the individual factors of country and page on conversion, we would now like to look at an interaction between page and country to see if there significant effects on conversion. Create the necessary additional columns, and fit the new model.

Provide the summary results, and your conclusions based on the results.

```
In [55]: df_cont = df_cont.rename(columns={'treatment': 'ab_page'})

In [59]: df_cont['ab_UK'] = df_cont['ab_page'] * df_cont['UK']
df_cont['ab_CA'] = df_cont['ab_page'] * df_cont['CA']
```

```

In [60]: df_cont['intercept'] = 1
         logistic_reg = sm.Logit(df_cont['converted'], df_cont[['intercept', 'ab_page', 'CA', 'UK'],

In [61]: results1= logistic_reg.fit()

Optimization terminated successfully.
      Current function value: 0.366109
      Iterations 6

In [62]: results1.summary2()

Out[62]: <class 'statsmodels.iolib.summary2.Summary'>
        """
                                Results: Logit
        =====
Model:                Logit                No. Iterations:    6.0000
Dependent Variable: converted                Pseudo R-squared: 0.000
Date:                2021-06-07 00:55 AIC:                212782.6602
No. Observations:    290584                BIC:                212846.1381
Df Model:            5                Log-Likelihood:    -1.0639e+05
Df Residuals:        290578                LL-Null:            -1.0639e+05
Converged:            1.0000                Scale:            1.0000
-----
                Coef.    Std.Err.    z        P>|z|    [0.025    0.975]
-----
intercept    -1.9865     0.0096   -206.3440   0.0000   -2.0053   -1.9676
ab_page      -0.0206     0.0137    -1.5052   0.1323   -0.0473    0.0062
CA           -0.0175     0.0377    -0.4652   0.6418   -0.0914    0.0563
UK           -0.0057     0.0188    -0.3057   0.7598   -0.0426    0.0311
ab_CA       -0.0469     0.0538    -0.8718   0.3833   -0.1523    0.0585
ab_UK        0.0314     0.0266     1.1807   0.2377   -0.0207    0.0835
=====
        """

```

0.2.1 Conclusions

-P-values has no significant change. So we accept the Null Hypotheses. There is no evidence that new page increase the conversion rate with comparing to the old one.

-Also the countries coefficient is very close (specially UK & US, Canada is slightly differ) as whole can't say that interaction to pages is very different from each others. There is no indication for the user choice from the different 3 countries.

-At the end I didn't see any benefits from new page. We didn't need to implement it, old page is doing well.

```

In [63]: from subprocess import call
         call(['python', '-m', 'nbconvert', 'Analyze_ab_test_results_notebook.ipynb'])

```

```
Out[63]: 0
```

```
In [ ]:
```